

Computer Vision



Lecture 2 Filtering and interest points

Lecturer:
Dr. Jan van Gemert (tud.computervision@gmail.com)

Recap of last time

- ▶ Course logistics
- ▶ Pinhole
- ▶ Reflection
- ▶ Questions?

Image Filtering



A bit of Image processing

How to get rid of noisy pixels?



Simple solution: replace pixel by neighbourhood average

Moving Neighborhood Average

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

Moving Neighborhood Average

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0								

Moving Neighborhood Average

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0								

Moving Neighborhood Average

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10							

Moving Neighborhood Average

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10							

Moving Neighborhood Average

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20						

Moving Neighborhood Average

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$G[x, y]$$

	0	10	20	30					

Moving Neighborhood Average

$$F[x, y]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

10 pixels

2 pixels lost to boundary (1 on each side)

$$G[x, y]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

8 pixels

How to prevent?


Smoothing by averaging



depicts box filter:
white = high value, black = low value




original



filtered

Averaging kernel


- Lets generalize to a kernel:
multiply weights per pixel and sum



- Which values in kernel H for the moving average example?

Averaging kernel

- Lets generalize to a kernel:
multiply weights per pixel and sum



- Which values in kernel H for the moving average example?

Correlation filtering

Say the averaging window size is $2k+1 \times 2k+1$:

$$G[i, j] = \underbrace{\frac{1}{(2k+1)^2}}_{\text{Attribute uniform weight to each pixel}} \underbrace{\sum_{u=-k}^k \sum_{v=-k}^k F[i+u, j+v]}_{\text{Loop over all pixels in neighborhood around image pixel } F[i,j]}$$

Attribute uniform weight to each pixel
Loop over all pixels in neighborhood around image pixel $F[i,j]$

Why $1/(2k+1)$?

Now generalize to allow different weights depending on neighboring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{Non-uniform weights}} F[i+u, j+v]$$

What happened to the $1/(2k+1)$?

Correlation filtering

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called cross-correlation, denoted $G = H \otimes F$

Filtering an image: replace each pixel with a linear combination of its neighbors.

The filter “kernel” or “mask” $H[u, v]$ is the prescription for the weights in the linear combination.

Questions?

Gaussian filter

- What if we want nearest neighboring pixels to have the most influence on the output?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$F[x, y]$


1	2	1
2	4	2
1	2	1

$\frac{1}{16}$


$H[u, v]$

This kernel is an approximation of a Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$




Smoothing with a Gaussian



The Gaussian is the only function that does not introduce artifacts
[koenderink, “The structure of images”, 1984]

Gaussian filters


- What parameters matter here? $h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$
- **Size** of kernel or mask
 - Note, Gaussian function has infinite support, but discrete filters use finite kernels




$\sigma = 5$ with 10
x 10 kernel

Gaussian filters

- What parameters matter here? $h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$
- **Size** of kernel or mask
 - Note, Gaussian function has infinite support, but discrete filters use finite kernels




$\sigma = 5$ with 10
x 10 kernel




$\sigma = 5$ with 30
x 30 kernel

Gaussian filters

- **Variance** of Gaussian: determines extent of smoothing



$\sigma = 2$ with 30
x 30 kernel



$\sigma = 5$ with 30
x 30 kernel

- Rule of thumb: set extent to 3σ

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0

?

Practice with linear filters



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

?

Practice with linear filters



Original

0	0	0
0	0	1
0	0	0



Shifted left
by 1 pixel
with
correlation

Practice with linear filters



Original

 $\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?


Practice with linear filters



Original

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1



Blur (with a
box filter)

Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

Convolution

- Convolution:
 - Flip the filter in both dimensions (bottom to top, right to left)
 - Then apply cross-correlation



How different from cross-correlation?

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$



*Notation for
convolution
operator*



Filtering an impulse signal

What is the result of convolving the impulse signal (image) F with the arbitrary kernel H ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$F[x, y]$

★

a	b	c
d	e	f
g	h	i

$H[u, v]$

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	a	b	c	0	0
0	0	d	e	f	0	0
0	0	g	h	i	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$G[x, y]$

Convolution vs. correlation

Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$

Cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

$$G = H \otimes F$$

For a Gaussian or box filter, how will the outputs differ?

If the input is an impulse signal, how will the outputs differ?

Properties of convolution

- Linear & shift invariant
- Commutative:

$$f * g = g * f$$

- Associative

$$(f * g) * h = f * (g * h)$$

- Identity:


$$\text{unit impulse } e = [\dots, 0, 0, 1, 0, 0, \dots]. \quad f * e = f$$

- Differentiation:

$$\frac{\partial}{\partial x}(f * g) = \frac{\partial f}{\partial x} * g$$

Questions?

Images are functions



- Edges look like steep cliffs

Images as functions


- We can think of an image as a function, f , from \mathbb{R}^2 to \mathbb{R} :
 - $f(x, y)$ gives the intensity at position (x, y)
 - How does reality differ from this definition?
 - Only defined over a rectangle, with finite range:

$$f: [a,b] \times [c,d] \rightarrow [0, 1.0]$$

- How about color?
 - “Vector valued” function
$$f(x, y) = \begin{bmatrix} r(x, y) \\ g(x, y) \\ b(x, y) \end{bmatrix}$$

Derivatives and edges

An edge is a place of rapid change in the image intensity function.




Effects of noise


Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal

$$f(x)$$




$$\frac{d}{dx}f(x)$$



Where is the edge?

Solution: smooth first



Where is the edge?


Look for peaks in $\frac{\partial}{\partial x}(h \star f)$

Derivative theorem of convolution


$$\frac{\partial}{\partial x}(h \star f) = \left(\frac{\partial}{\partial x}h\right) \star f$$

Differentiation property of convolution.


f




$\frac{\partial}{\partial x}h$





$\left(\frac{\partial}{\partial x}h\right) \star f$



Derivative of Gaussian filters



x-direction



y-direction







Image gradient

The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity



The gradient direction (orientation of edge normal θ) is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Example




original image (Lena)

Compute Gradients




Derivative of Gaussian



Derivative of Gaussian

- Which one is the gradient in the x-direction (resp. y-direction)?

Gradient Orientation




Orientation of the gradient

$$\text{theta} = \text{atan2}(\text{gy}, \text{gx})$$


Gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$




- Do you see a problem here?


Edge detection, Take 2

 f


Edge

 $\frac{d^2}{dx^2} g$


Second derivative
of Gaussian
(Laplacian)

 $f * \frac{d^2}{dx^2} g$


Edge = zero crossing
of second derivative

Questions?


- ▶ Next topic: interest point detection

Applications


- Feature points (blobs) are used for:
 - Image alignment
 - 3D reconstruction
 - Motion tracking
 - Robot navigation
 - Indexing and database retrieval
 - Object recognition




Edge detection, Take 2

 f


Edge

 $\frac{d^2}{dx^2} g$



Second derivative
of Gaussian
(Laplacian)

 $f * \frac{d^2}{dx^2} g$


Edge = zero crossing
of second derivative

From edges to blobs


- Edge = ripple
- Blob = superposition of two ripples



Spatial selection: the magnitude of the Laplacian response will achieve a maximum at the center of the blob, provided the scale of the Laplacian is “matched” to the scale of the blob

Scale selection


- We want to find the characteristic scale of the blob by convolving it with Laplacians at several scales and looking for the maximum response
- However, Laplacian response decays as scale increases:



Why does this happen?

Scale normalization

- The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases




Scale normalization






- The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases
- To keep response the same (scale-invariant), must multiply Gaussian derivative by σ
- Laplacian is the second Gaussian derivative, so it must be multiplied by σ^2

Effect of scale normalization






Original signal



Unnormalized Laplacian response





Scale-normalized Laplacian response



↑
maximum

Blob detection in 2D



- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D



$$\nabla^2 g = \frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2}$$

Blob detection in 2D


- Laplacian of Gaussian: Circularly symmetric operator for blob detection in 2D



Scale-normalized:
$$\nabla_{\text{norm}}^2 g = \sigma^2 \left(\frac{\partial^2 g}{\partial x^2} + \frac{\partial^2 g}{\partial y^2} \right)$$

Characteristic scale


- We define the characteristic scale of a blob as the scale that produces peak of Laplacian response in the blob center



T. Lindeberg (1998). ["Feature detection with automatic scale selection."](#) *International Journal of Computer Vision* **30** (2): pp 77--116.

Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales
2. Find maxima of squared Laplacian response in scale-space



Scale-space blob detector: Example




Scale-space blob detector: Example




sigma = 11.9912

Scale-space blob detector: Example




Questions?

A feature detector example:
Harris corner detector




The Basic Idea


- We should easily localize the point by looking through a small window
- Shifting a window in *any direction* should give *a large change* in intensity




Harris Detector: Basic Idea



“flat” region:
no change as shift
window in all
directions



“edge”:
no change as shift
window along the
edge direction



“corner”:
significant change as
shift window in all
directions

Harris Detector: Mathematics

Window-averaged change of intensity induced by shifting the image data by $[u, v]$:


$$E(u, v) = \sum_{x, y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

Window
function


Shifted
intensity

Intensity

Window function $w(x, y) =$



1 in window, 0 outside



Gaussian

Taylor series approx to shifted image

$$E(u, v) = \sum_{x, y} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

$$E(u, v) \approx \sum_{x, y} w(x, y) [I(x, y) + uI_x + vI_y - I(x, y)]^2$$

$$= \sum_{x, y} w(x, y) [uI_x + vI_y]^2$$

$$= \sum_{x, y} w(x, y) \begin{pmatrix} u & v \end{pmatrix} \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix} \begin{pmatrix} u \\ v \end{pmatrix}$$

Harris Detector: Mathematics

Expanding $I(x,y)$ in a Taylor series expansion, we have, for small shifts $[u, v]$, a *bilinear* approximation:

$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

M is also called “structure tensor”
or “second moment matrix”

Harris Detector: Mathematics

Intensity change in shifting window:


$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

Where is the largest change?



λ_1, λ_2 , eigenvalues of M

Ellipse $E(u, v) = \text{const}$

Iso-intensity contour of $E(u, v)$





Selecting Good Features



How large is each of the two eigenvalues?

small λ_1 , small λ_2



Selecting Good Features



How large is each of the two eigenvalues?

large λ_1 , small λ_2

Selecting Good Features




How large is each of the two eigenvalues?

λ_1 and λ_2 are large

Harris Detector: Mathematics

Classification of image points using eigenvalues of M :




Harris Detector: Mathematics

No need to compute eigenvalues:

$$\det M = \lambda_1 \lambda_2$$

$$\text{trace } M = \lambda_1 + \lambda_2$$

λ_1 and λ_2 are small;
 E is almost
 constant in all
 directions



Harris Detector: Mathematics

No need to compute eigenvalues:

$$\det M = \lambda_1 \lambda_2$$


$$\text{trace } M = \lambda_1 + \lambda_2$$

Measure of corner response:

$$R = \det M - k (\text{trace } M)^2$$

(k – empirical constant, $k = 0.04$ - 0.06)

- R depends only on eigenvalues of M
- R is large for a **corner**
- R is negative with large magnitude for an **edge**
- $|R|$ is small for a **flat** region



Harris Detector


- The Algorithm:
 - Compute each element of the structure tensor on the image level
 - Find points with large corner response function R ($R > \text{threshold}$)
 - Take the points of local maxima of R

Harris Detector: Steps




Harris Detector: Steps

Compute corner response R



Harris Detector: Steps

Find points with large corner response: $R > \text{threshold}$



Harris Detector: Steps

Take only the points of local maxima of R



Harris Detector: Steps



Invariance and covariance

- We want features to be *invariant* to photometric transformations and *covariant* to geometric transformations
 - **Invariance:** image is transformed and features do not change
 - **Covariance:** if we have two transformed versions of the same image, features should be detected in corresponding locations



Harris Detector: Summary

- Average intensity change in direction $[u, v]$ can be expressed as a bilinear form:

$$E(u, v) \cong [u, v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

- Describe a point in terms of eigenvalues of M :
measure of corner response

$$R = \lambda_1 \lambda_2 - k (\lambda_1 + \lambda_2)^2$$


- A good (corner) point should have a *large intensity change in all directions*, i.e. R should be large positive

Ideal feature detector

- Would always find the same point on an object, regardless of changes to the image.
- I.e., insensitive to changes in:
 - Scale
 - Lighting
 - Perspective imaging
 - Partial occlusion


Harris Detector: Some Properties

- Rotation covariance?



Harris Detector: Some Properties

- Rotation invariance?



Ellipse rotates but its shape (i.e. eigenvalues) remains the same



Corner response R is invariant to image rotation

Harris Detector: Some Properties

- Invariance to image intensity change?

Harris Detector: Some Properties

- Invariance to image intensity change?
- Partial invariance to additive and multiplicative intensity changes
 - ✓ Only derivatives are used \Rightarrow invariance to intensity shift $I \rightarrow I + b$
 - ✓ Intensity scale: $I \rightarrow a I$ Because of fixed intensity threshold on local maxima, only partial invariance to multiplicative intensity changes.




Harris Detector: Some Properties

- Invariant to image scale?

Harris Detector: Some Properties

- Not invariant to *image scale*!



All points will be
classified as **edges**

Corner !

Questions?

- Scale invariant Harris detector?

Summary

- Filtering
- Convolution and cross-correlation
- Gaussian filtering
- Gradients
- Laplacian for blob detection
- Laplacian for scale selection
- Harris detector